

ROS 2

Advanced communication II

Roberto Masocco

`roberto.masocco@uniroma2.it`

University of Rome Tor Vergata

Department of Civil Engineering and Computer Science Engineering

May 16, 2024



TOR VERGATA
UNIVERSITY OF ROME

School of Engineering

Recap

Messages are the most basic, **one-way** communication paradigm.

QoS policies and other topic settings affect communication behavior among entities.

Services are a simple implementation of a **two-way**, **client-server** communication paradigm.

This lecture is [here](#).

Updates

- **Updated code examples**, please pull new commits!
- Follow-up on **message topics code examples**:
 - ▶ [resetting_sub](#) example.

Roadmap

1 Interface packages

2 Actions

1 Interface packages

2 Actions

Defining custom interfaces

From package organization to build

Custom communication interfaces can be defined in **appropriate packages**.

The organization of such packages is regulated by a set of **best practices**.

When these packages are built, they generate both **code to include in** and **libraries to link against** the packages using them.

For everything to work, each package must have **the same interface packages** as **dependencies**, as well as link against **the same generated code**.

Find all relevant information in our [interfaces.md](#) file!

Roadmap

1 Interface packages

2 Actions

Limitations of services

The third paradigm exists because services rely on the following **restrictive assumptions**.

Services implementation assumptions

- Since the client may block for the entire duration of the request processing, **server computations should be short and always produce some result** (e.g., even an error must be a result, but **we** have to encode it).
- Service calls are finished only when the response has been received, *i.e.*, **if the server crashes before sending the response, the client's behavior is undefined, especially when it is blocking** (no **state machine**! Say hello to **deadlocks**, crashes...).
- Once a service is called, **the processing of the request may never be interrupted**.

These make operations that **must be requested** and **take a long time** (for CPUs!) completely unfeasible.

Think of real stuff such as **movement**, **navigation**...

ROS 2 actions

Full client-server paradigm

Built on services and message topics, they **decouple computations from middleware APIs**, thanks to three concepts that embody the **three stages of the communication**:

- ➊ **Goal**: the full request of the operation to be executed.
- ➋ **Feedback**: intermediate results and information about the ongoing processing.
- ➌ **Result**: the final result of the requested operation.

Their implementation is still a bit cumbersome because of the **many different data types** (classes) involved, as well as many **API inconsistencies**.

It is found in the [rclcpp_action](#) and [rclpy action](#) libraries.

They are **extensively used to implement robot navigation and movement**, and many more complex real-world operations, providing a **software abstraction** for them.

ROS 2 actions

Full client-server paradigm




Figure 1: Example of an action server and *client*.

ROS 2 actions

The goal state machine

Goal State Machine





Figure 2: State machine¹ of an action goal, implemented and managed internally by ROS 2.

¹[Actions - ROS 2 Design](#)

ROS 2 actions

Communication overview

In actual ROS 2 applications, the **client** requests the completion of some **goal** to the **server**. The middleware only offers APIs to **notify the state of the goal** between the two.

- ➊ The **client** sends a **goal service request** to the server.
- ➋ The **server** may **accept** or **reject** the goal request.
- ➌ Server computations are usually started when the goal is **executed**: the middleware only keeps track the state of the goal, its updates and the rest are up to the developer.
- ➍ The **client may cancel** the goal request; the **server may abort** the goal request; intermediate results and information, if any, are published by the server on the **feedback topic**.
- ➎ The **client** asks the server for the final result over the **result service**.

ROS 2 actions

CLI introspection tools

The main command is `ros2 action` with the following verbs:

- `list` Lists all active actions.
- `info` Prints information about an action.
- `send_goal` Sends a goal request to an action server, and prints the result; with `-f` prints also feedback messages.

ROS 2 actions

Coding hints for servers and clients

Servers

Goal requests are handled with **callbacks**, while computations can be handled freely (usually done in **separate threads**).

Cancellation requests are handled with **callbacks** and can be **polled** during the computation. When done, the goal must be marked as **succeeded** or **aborted**.

Clients

Similarly to services, much is done with **future objects**, but **callbacks may be defined** to handle **goal**, **result** and **cancellation responses**, and **feedbacks**.

Handling all possible scenarios for a goal results in the **longest and most complicated code that a ROS 2 application may ever require.** 😊

Interface files

Actions

Combine **three messages** in a single interface file, separated by ---.

Action file names end with `.action`.

```
1 # GOAL
2 int32 order
3 ---
4 # RESULT
5 int32[] sequence
6 ---
7 # FEEDBACK
8 int32[] partial_sequence
```

Listing 1: Definition of the `ros2_examples_interfaces/action/Fibonacci` action.

Example

Fibonacci computer

Now go have a look at the [ros2-examples/src/cpp/actions_example_cpp](#) package!

If you're curious, the [ros2-examples/src/cpp/advanced/complete_actions_cpp](#) package, which implements the complete goal state machine using a multithreaded executor.

- Run the action client and server examples, and try to call the action from the command line.
- Modify the feedback message: instead of a partial sequence, it should publish the length of the sequence so far; this requires:
 - ① modifying the action definition in `ros2_examples_interfaces`;
 - ② modifying the server node to publish the length of the sequence instead of the partial sequence in the feedback message (hint: use methods of the `std::vector` class to get the length of the partial sequence in one go);
 - ③ modifying the feedback callback in the client to parse and print the length from the feedback message.